

Advanced Topics in Communication Networks

Yannick Merkli, ymerkli@ethz.ch

ETH Zurich, HS 2019

Abstract

This document is a "lecture summary" style script that closely follows the slides of the *Advanced Topics in Communication Networks* lecture at ETH Zurich [8]. The contribution to this is editing and refactoring as well as providing additional material for better understanding. This summary was created during the fall semester 2019. Due to updates to the syllabus content, some material may no longer be relevant for future versions of the lecture. Most graphics are copy & pasted from the slides. If you don't want yours here, please contact me and I will remove them. Otherwise, this work is published as CC BY-NC-SA.



I do not guarantee correctness or completeness, nor is this document endorsed by the lecturers. Feel free to point out any erratas. For the full L^AT_EX source code, consider github.com/ymerkli/eth-summaries.

Contents

1	Introduction	1
1.1	The network management crisis	1
1.2	Software-defined networking (SDN)	1
1.3	Deep network programmability	1
2	The P4 programing language	2
2.1	Language specification	3
2.2	Parser	4
2.3	Match-action tables	4
2.4	Actions	4
2.5	Control flow	5
2.6	Stateful objects	5
2.7	Hash functions	6
3	Probabilistic data structures	7
3.1	Bloom filter	7
3.2	Counting bloom filter	8
3.3	Invertible Bloom Lookup Tables (IBLT)	9
3.4	Sketches	10
4	P4 hardware targets	12
4.1	Parsing	13
4.2	Logical pipeline	13
4.3	Barefoot Tofino	14
4.4	P4 switch design	17
5	P4-based applications	17
5.1	CONGA: Distributed Congestion-Aware Load Balancing for Datacenters	17
5.2	INT: In-band Network Telemetry	17
5.3	NetCache: Balancing Key-Value Stores with Fast In-Network Caching	18
5.4	p4v: Practical Verification for Programmable Data Planes	18
5.5	Others	18
6	Exercise design summaries	19
6.1	03-ECMP	19
6.2	03-Flowlet-switching	19
6.3	04-CountMin-sketch	20
6.4	04-Heavy-hitter-detector	21
6.5	05-Simple-routing	21
6.6	05-Traceroutable	22
6.7	06-Packet-Loss-Detection	23
7	Oral exam questions	25
	References	30

1 Introduction

Networking is on the verge of a paradigm shift towards deep programmability.

1.1 The network management crisis

Networks are large distributed systems running a set of distributed algorithms. These algorithms produce the forwarding state which drives IP traffic to its destination. Operators adapt their network behavior by configuring each network device individually. This is extremely tedious and error-prone with a single mistyped line being enough to bring down an entire network (fat-thumbings). Further, the complexity in networks keeps increasing with more and more protocols appearing, a lot of which are badly documented (read an RFC and find out).

1.2 Software-defined networking (SDN)

SDN tries to design network control and is predicated around two simple concepts: (1) Separate the control-plane from the data-plane. (2) Provide an API to directly access the data-plane. In traditional computer networks, each networked device has a local control-plane. SDN allows to have a central control-plane, controlling multiple networked devices at once. This has several advantages: (1) Simpler management, (2) Faster pace of innovation, (3) Easier interoperability, (4) Simpler, cheaper equipment. Having a common open, vendor-agnostic interface enables a control plane to control forwarding devices from different hardware and software vendors. OpenFlow does exactly this: OpenFlow is essentially an API to a switch flow table. The OpenFlow interface started simple, with the abstraction of a single table of rules that could match packets on a dozen header fields (e.g., MAC addresses, IP addresses, protocol, TCP/UDP port numbers, etc.). Over the past five years, the specification has grown increasingly more complicated, with many more header fields and multiple stages of rule tables, to allow switches to expose more of their capabilities to the controller. So essentially, the OpenFlow protocol became too complex.

1.3 Deep network programmability

Deep network programmability tries to adopt the good ideas of OpenFlow while solving its shortcomings. OpenFlow's problem is that it's not flexible enough for a highly dynamic environment such as networking, with constantly changing protocols and specifications. Future switches should support flexible mechanisms for parsing packets and matching header fields, allowing controller applications to leverage these capabilities through a common, open interface. Recent chip designs show that such flexibility can be achieved in custom ASICs at terabit speeds.




Figure 1: Protocol Independent Switch Architecture (PISA) for high-speed programmable packet forwarding. [8]

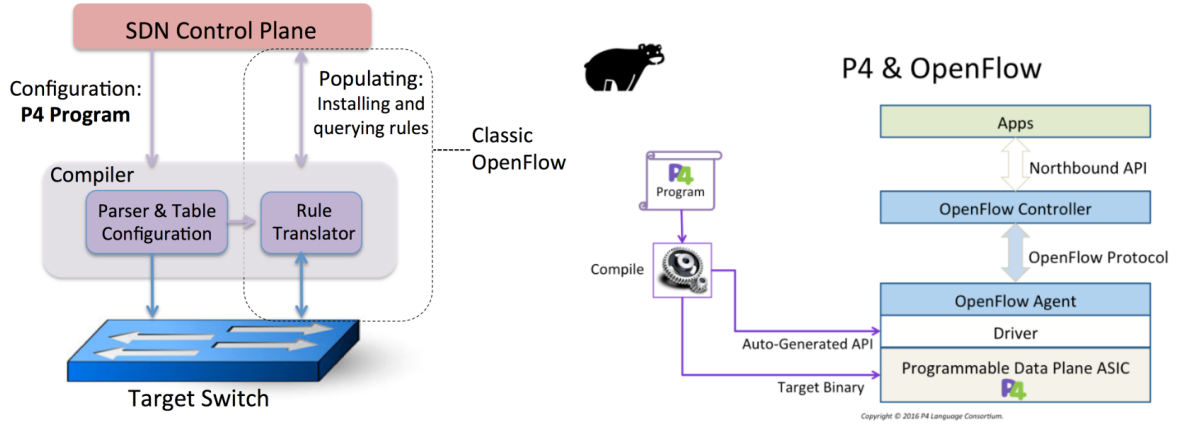


Figure 2: P4 is a language to configure switches. [3]

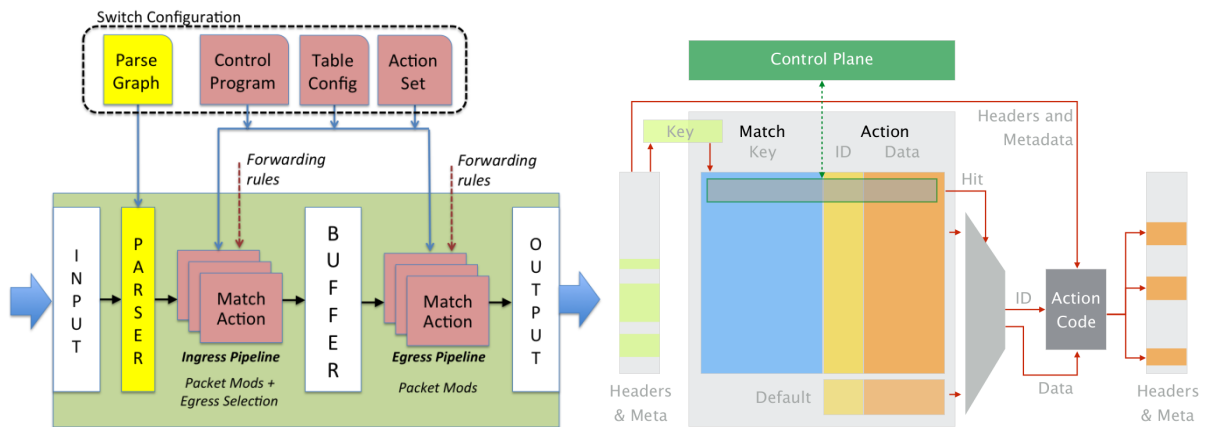
Each chip has its own low-level interface, however this might vary for different hardware. Ideally we would want a higher-level language which can be used to configure a switch. This is exactly what P4 (Programming Protocol-independent Packet Processors) does: P4 is a higher-level language which is used to configure a switch, telling it how packets are to be processed. P4 can further be used with existing APIs (such as OpenFlow) that are designed to populate the forwarding tables in fixed function switches.

2 The P4 programming language

P4 raises the level of abstraction for programming the network, and can serve as a general interface between the controller and the switches. As such, P4 tries to achieve the following three main goals:

- **Reconfigurability.** The controller should be able to redefine the packet parsing and processing in the field.
- **Protocol independence.** The switch should be able to specify (i) a packet parser (ii) a collection of match-action tables that process these headers.
- **Target independence.** The P4 program should run on various hardware with the compiler producing the target-dependent program.

A P4 program consists of three basic parts: Parser, match-action pipeline, deparser. In this course, we rely on a simple $P4_{16}$ switch architecture (v1model).



(a) The abstract P4 forwarding model. [3]

(b) Reconfigurable match-action table. [8]

The forwarding model is controlled by two types of operations: Configure and Populate. Configure operations program the parser, set the order of match + action stages, and specify the header fields processed by each stage. Configuration determines which protocols are supported and how the switch may process packets. Populate operations add (and remove) entries to the match-action tables that were specified during configuration. Population determines the policy applied to packets at any given time.

2.1 Language specification

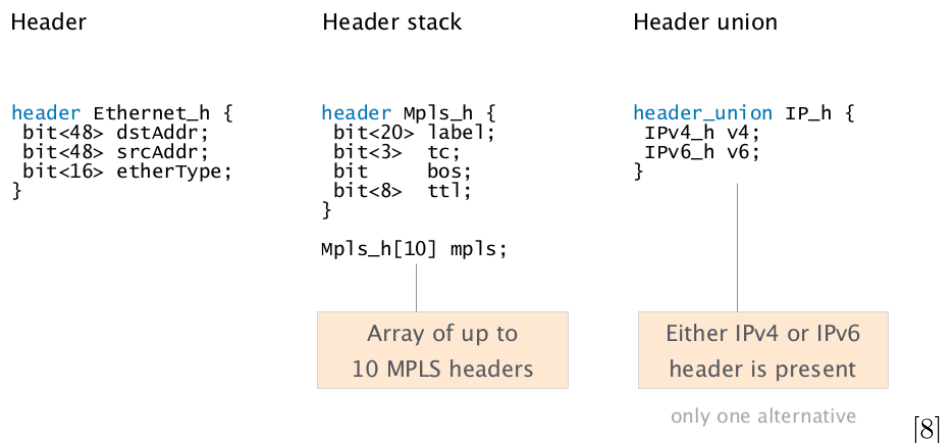
2.1.1 Data types

$P4_{16}$ is a statically typed language with base types and operators to derive composed ones. Base types are:

bool	Boolean value
bit<W>	Bit-string of width W
int<W>	Signed integer of width W
varbit<W>	Bit-string of dynamic length $\leq W$
match_kind	describes ways to match table keys
error	used to signal errors
void	no values, used on few restricted instances

Note that there are no floats or strings.

Headers are composed operators. Headers are similar to structs in C, containing different fields. Parsing a packet using `extract()` fills in the fields of the header from a network packet. Headers further have a hidden "validity" field which is set to true upon a successful `extract()`.



2.1.2 Operations

P4 operations are similar to C operations and vary depending on the types (unsigned/signed ints, ...). However there is no division or modulo.

Constants, variable declarations and instantiations are pretty much the same as in C too. However, Variables have local scope and their values is not maintained across subsequent invocations. This is due to the fact that the code will be rerun for every new packet. We thus cannot use variables as states since they will be erased for each subsequent run of the code. In order to maintain state, you have to use tables or extern objects.

2.1.3 Statements

P4 statements are pretty classical too:

- `return`
- `exit`
- `if(){...} else {...}` (not in parsers)
- `switch (t.apply.action_run) { a1: {...} a2: {...}}` (only in control blocks)

Loops do not exist in P4 (with one exception, see 2.2)

2.2 Parser

The parser uses a state machine to map packets into headers and metadata. Parsing a header stack requires the parser to loop. This is the only 'loops' that are possible in P4.

Defining (and parsing) *custom* headers allow you to implement your own protocol. This is why OpenFlow didn't work well in reality and P4 does: with P4, companies can just define their own protocols (for example ETH and UZH can establish a common tunneling protocol). With OpenFlow, companies were either bound to existing protocols and headers or, in case of a influential company, they could bring up new protocols which were eventually implemented in OpenFlow and available to everyone, even though only a small subset of people need that protocol. This led to OpenFlow becoming the overblown beast it is today.

2.3 Match-action tables

Match-action tables are the mechanism for performing packet processing. The P4 program defines the fields on which a table may match and the actions it may execute.

Tables can match on one or multiple keys in different ways:

- **exact**: exact comparison (0x01020304)
- **ternary**: compare with mask (0x01020304 & 0x0F0F0F0F)
- **lpm**: longest prefix match (0x01020304/24)
- **range**: check if in range (0x01020304 0x010203FF)

Table entries are added through the control plane.

Example: `table_add ipv4_lpm ipv4_forward 1.2.3.0/24 => 01:01:01:01:01:01 1`

The match-action tables are divided between ingress and egress. While both may modify the packet header, ingress match-action determines the egress port(s) and determines the queue into which the packet is placed. Based on ingress processing, the packet may be forwarded, replicated (for multicast, span, or to the control plane), dropped, or trigger flow control. Egress match-action performs per-instance modifications to the packet header e.g., for multicast copies.

Packets can carry additional information between stages, called metadata, which is treated identically to packet header fields. Some examples of metadata include the ingress port, the transmit destination and queue, a timestamp that can be used for packet scheduling, and data passed from table-to-table that does not involve changing the parsed representation of the packet such as a virtual network identifier.

2.4 Actions

Actions are blocks of statements that possibly modify the packets (think of them as functions in C). P4 supports the construction of complex actions from simpler protocol-independent primitives. Actions can either be invoked from within a control block or as a result of a match in a match-*action*-table.

Actions that are invoked from within a control block take directional parameters, indicating how the corresponding value is treated within the block.

- **in**: parameters in only read inside the action (like parameters to a function)
- **out**: parameter is uninitialized and will be written to inside the action (like return values)
- **inout**: combination of in and out (like call by reference)

Action parameters resulting from a table lookup do not take a direction as they come from the control plane.

2.5 Control flow

Control flow consists of various concepts, for an extensive list, please consider [9]. Some basic concepts are:

- applying a table: `ipv4_lpm.apply()`
- Checking if there was a hit: `if (ipv4_lpm.apply().hit) {...}`
- Check which action was executed: `switch (ipv4_lpm.apply().action_run) {...}`

Other often used concepts are:

- re-computing checksums (needs to be done when changing a packet, otherwise the kernel will drop packets upon checksum mismatch)
- cloning packets
- sending packets to control plane (using dedicated Ethernet port, or target-specific mechanisms (e.g. digests))

2.6 Stateful objects

As mentioned in section 2.1.2, variables have local scope and their value is not maintained across subsequent invocations. In order to build stateful apps, we thus need stateful objects. Match-action tables are one way of achieving persistent state, however the P4 *v1model* offers more stateful objects as externs.

2.6.1 Registers

Registers are assigned in arrays and are useful for storing (small amounts of) arbitrary data. We can read from and write to a specified register index in an array of registers. Further, registers can be read from the control plane. The syntax works as follows:

```
register<bit<48>>(16384) last_seen; bit<48> last_pkt_ts;
last_seen.read(last_pkt_ts, flow_id);
last_seen.write(flow_id, standard_metadata.ingress_global_timestamp);
```

Registers can be used to build a stateful firewall, for example.


2.6.2 Counters

Counters are for counting (unsurprisingly). A counter can be increased from the data plane *but* can only be read from the control plane. Like registers, counters are assigned in arrays and an index needs to be specified when executing a `count()` operation. Counters can be one of three different types: packets, bytes, packet_and_bytes (each type counts what it specifies to count). The syntax works as follows:

```
counter(512, CounterType.packets_and_bytes) port_counter;
port_counter.count((bit<32>)standard_metadata.ingress_port);
```

Reading from the control plane:

```
RuntimeCmd: counter_read MyIngress.port_counter 1
MyIngress.port_counter[1]= BmCounterValue(packets=13, bytes=1150)
```



Direct counters are a specially kind of counters that are attached to tables. We can specify a counter in the table definition and thus attach it to the table. Then, each table entry has a counter cell that counts when the entry matches.

```
direct_counter(CounterType.packets_and_bytes) direct_port_counter;
table count_table {
    ...
    counters = direct_port_counters;
}
```

2.6.3 Meters

Meters are used for rate-limiting. Meters 'color' flows into green, yellow and red (see two-rate three-color meters, RFC2698). We have two threshold rates: the committed information rate (CIR) and the peak information rate (PIR). Green flows neither exceed the CIR nor the PIR. Yellow flows exceed the CIR but don't exceed the PIR. Red flows exceed the PIR. Like counters, meters can be one of three different types: packets, bytes, packet_and_bytes and they are assigned in arrays. And just like for counters, we also have direct meters which can be attached to tables. The syntax looks as follows:

```
meter(32w16384, MeterType.packets) my_meter;
my_meter.execute_meter<bit<32>>(meter_index, meta.meter_tag); //action
```

Executing the meter will yield one of 3 values: 0 (GREEN), 1 (YELLOW) or 2 (RED). This value will be copied to metadata field meta.meter_tag

2.7 Hash functions

P4 *v1model* offers hash functions:

```
enum HashAlgorithm{crc32,crc32_custom,crc16,s,random,identity,csum16,xor16}
extern void hash<O,T,D,M>(out O result, in HashAlgorithm algo, in T base, in D data,
in M max);
```

Object	Data plane interface		Control plane interface	
	read	modify/write	read	modify/write
Table	apply()	—	yes	yes
Register	read()	write()	yes	yes
Counter	—	count()	yes	reset
Meter	execute()		configuration only	

Figure 5: Stateful objects summary. [8]


3 Probabilistic data structures

P4 provides us with built-in stateful data structures such as arrays of registers, counters or meters. However, with these we need to deal with severe limitations such as limited number of operations and memory. That's why we consider more advanced stateful data structures.

3.1 Bloom filter

We want a data structure for insertion and membership queries. In order to get a deterministic number of required operations, the data structure should be probabilistic. A simple approach would be a basic fixed size table with M 1-bit cells for N elements. If we observe an element, we hash it and set the bit at the corresponding index to 1. Due to hash collision ($N > M$), we get a false positive rate of $FPR = 1 - (1 - \frac{1}{M})^N$. The advantage of this approach is that only one operation is required per insertion or query. However, roughly 100x more cells are required than the number of element we want to store for a 1% false positive rate.

Bloom filters are a more memory-efficient approach for insertions and membership queries. Bloom filters consist of a fixed size table **bf** with M 1-bit cells and K hash functions and we write a 1 at each position indicate by each hash function.



- Insert e into **bf**:
 1. $\forall i \in [1, K]$, calculate $h_i(e)$
 2. $\mathbf{bf}[h_i(e)] = 1, \forall i \in [1, K]$
- Membership query e :
 1. if $\mathbf{bf}[h_i(e)] == 1, \forall i \in [1, K]$
 $\rightarrow e$ is in **bf**
 2. else
 $\rightarrow e$ is not in **bf**

Figure 6: Bloom filter. [8]

The advantage of bloom filters is that they use about 10x less memory than the simple approach. However, they require slightly more operations.

3.1.1 Dimension your bloom filter

N elements, M cells, K hash functions, FP false positive rate.

- probability that one hash function returns the index of a particular cell: $\frac{1}{M}$
- probability that one hash function does not return the index of a particular cell: $1 - \frac{1}{M}$
- probability of a cell to be 0: $(1 - \frac{1}{M})^{KN}$
- false positive rate $P(FP)$: $(1 - (1 - \frac{1}{M})^{KN})^K$
- false negative rate: 0

For an approximation, use: $p := P(FP) = (1 - (1 - \frac{1}{M})^{KN})^K \cong (1 - e^{-KN/M})^K$.

$FN = 0$ since if we have seen an element once, all cells corresponding to its hashes will have been set to 1. Since we never delete elements, $FN = 0$. If we were to delete elements from a bloom filter (i.e. set all cells corresponding to an element to 0), false negatives would be possible. Because deletions are not possible, the controller may need to regularly reset the bloom filter. Resetting a bloom filter takes some time during which it is not usable. A common trick is to use two bloom filters and use one when the controller resets the other one.

We can't just keep increasing K to lower FP since for too many hash functions, we will have lots of 1s in the table and a lower chance of finding 0s, which increases FP .

There's a global minimum when $K = \ln(2) * \frac{M}{N} \approx 0.7 * \frac{M}{N}$, found by taking derivative of $P(FP)$. For that choice of K , resulting $p := P(FP) = 2^{-K} \approx 0.6185^{M/N}$.

Given optimal K , choice of optimal $M = -\frac{N \ln p}{(\ln 2)^2} \rightarrow O(N)$ space.

```
control MyIngress(...) {
    register register<bit<1>>(NB_CELLS) bloom_filter;

    apply {
        hash(meta.index1, HashAlgorithm.my_hash1, 0,
            {meta.dstPrefix, packet.ip.srcIP}, NB_CELLS);
        hash(meta.index2, HashAlgorithm.my_hash2, 0,
            {meta.dstPrefix, packet.ip.srcIP}, NB_CELLS);

        if (meta.to_insert == 1) {
            bloom_filter.write(meta.index1, 1);
            bloom_filter.write(meta.index2, 1);
        }

        if (meta.to_query == 1) {
            bloom_filter.read(meta.query1, meta.index1);
            bloom_filter.read(meta.query2, meta.index2);


            if (meta.query1 == 0 || meta.query2 == 0) {
                meta.is_stored = 0;
            }
            else {
                meta.is_stored = 1;
            }
        }
    }
}
```

Figure 7: Bloom filter with $K = 2$ hash functions implemented in $P4_{16}$. [8]

3.2 Counting bloom filter

Since we cannot delete items from a bloom filter, we extend them to handle deletions. For counting bloom filters, to add an element, increment the corresponding counters by 1. The cells are thus no longer 1-bit but larger. To delete an element, decrement the corresponding counters by 1. An element is considered in the table if all counters are larger than 0. All of our prior analysis for standard bloom filters applies to counting bloom filters. We still have $FN = 0$ since a counter of a cell only reaches 0 if all elements mapped to that cell were deleted. Unless we have counter overflow. If a counter gets too large and wraps around to 0, false negatives are possible. Counters must be large enough to avoid overflow. Poisson approximation suggests 4 bits/counter.

Thus, Counting Bloom Filters do handle deletions at the price of using more memory



- Insert e into cbf:
 1. $\forall i \in [1, K]$, calculate $h_i(e)$
 2. $\text{cbf}[h_i(e)] += 1, \forall i \in [1, K]$
- Membership query e :
 1. if $\text{bf}[h_i(e)] \geq 1, \forall i \in [1, K]$
 $\rightarrow e$ is in cbf
 2. else
 $\rightarrow e$ is not in cbf

Figure 8: Counting bloom filter. [8]

3.3 Invertible Bloom Lookup Tables (IBLT)

Invertible Bloom Lookup Tables (IBLT) stores key-value pairs and allows for lookups and a complete listing. Each cell contains three fields:

- count: counts the number of entries mapped to this cell
- keySum: the sum of all keys mapped to this cell
- valueSum: the sum of all the values mapped to this cell

In many settings, we can use XORs in place of sums, for example to avoid overflow issues.

Add a new key-value pair	Delete a key-value pair (assuming it's in set)
for each hash function hash the key to find the index then at this index increment the count by one add key to keySum add value to valueSum	for each hash function hash the key to find the index then at this index subtract one from the count subtract key from keySum subtract value from valueSum

The value of a key can be found if the key is associated to **at least** one cell without a hash collision, i.e. a cell with *count* = 1 (see 9b).

Listing the IBLT
While there is an index for which count = 1 Find the corresponding key-value pair and return it Delete the corresponding key-value pair

A complete listing is not possible if there is a key that has *count* > 1 at each cell indicated by the hash of the key. Unless the number of iterations is very low, loops cant be implemented in hardware. The listing is done by the controller.

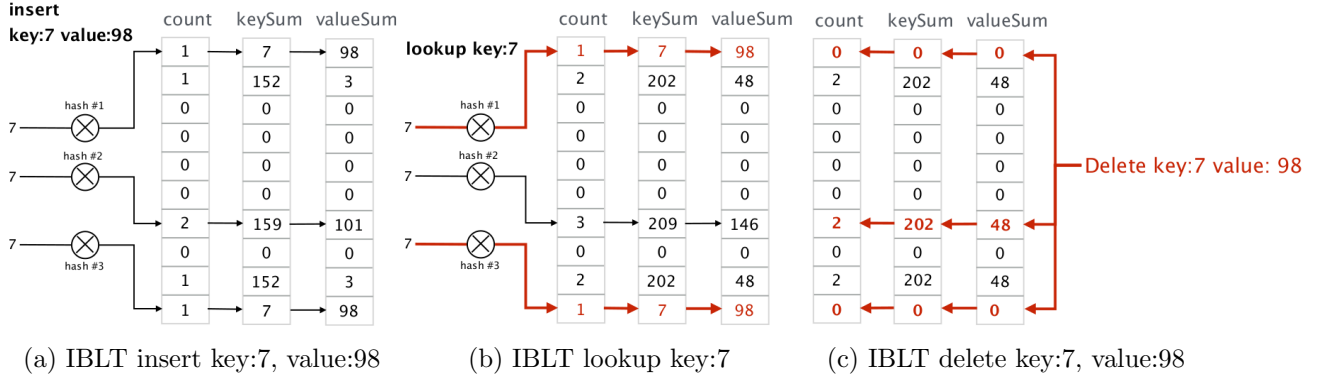


Figure 9: Invertible Bloom Lookup Tables (IBLT) insertion and lookup [8]

3.4 Sketches

Bloom filters allow us to do efficient insertion and membership queries at the cost of false positives. This way we can quickly filter only those elements that might be in the set. However, this is often not enough for monitoring use cases. We also want to approximate frequencies of elements in a data stream. Again, we use probabilistic measures which make this more efficient by allowing miscounting.

3.4.1 CountMin sketch

A CountMin sketch uses the same principles as a counting bloom filter, but is designed to have provable L1 error bounds for frequency queries. We use the following notation: a vector of frequencies (counts) of all distinct elements x_i : $\vec{x} = [x_1, x_2, \dots]^T$.

$$Pr[\hat{x}_i - x_i \geq \epsilon \|\mathbf{x}\|_1] \leq \delta$$

with \hat{x}_i : estimated frequency, x_i : true frequency, $\|\mathbf{x}\|_1$: sum of frequencies.

A CountMin Sketch uses multiple arrays and hashes. Let d be the number of arrays, with one hash function per array and w indices per array (range of hashes). We thus have $w * d$ counters. CountMin sketches are essentially Bloom filters where we increase the counter at each position indicated by the hashes. Hash collisions can cause over-counting.

Count a new element	Lookup the count of an element
for each hash function	for each hash function
hash the element to find the index	hash the element to find the index
then at this index	then at this index
increment the count by one	lookup the count
	return the min count

CountMin sketch recipe

Choose $d = \left\lceil \ln \frac{1}{\delta} \right\rceil$, $w = \left\lceil \frac{e}{\epsilon} \right\rceil$

Then $\hat{x}_i - x_i \geq \epsilon \|\mathbf{x}\|_1$ with a probability less than δ

Figure 10: CountMin sketch recipe, with e : Euler's number [8]




Figure 11: CountMin sketches count and lookup [8]

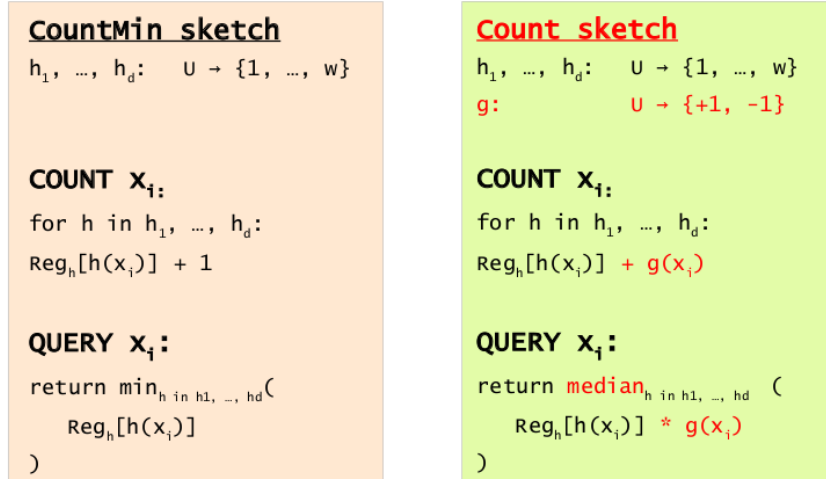


Figure 12: CountMin sketch and Count sketch [8]

3.4.2 Count sketch

A Count sketch uses the same principles as a counting bloom filter, but is designed to have provable L2 error bounds for frequency queries. The Count sketch uses additional hashing to give L2 error bounds, but requires more resources.


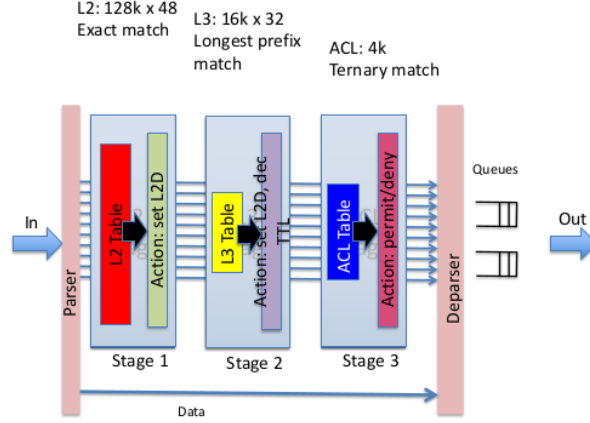


Figure 13: Count sketch recipe [8]

4 P4 hardware targets


How can we allow network programmability in the field, at reasonable cost, and without sacrificing speed. Let's look at a concrete design: Reconfigurable Match Tables (RMT) [4]. This paper argues that flexibility does not come at the price of performance or cost.

Let's first look at a fixed-function switch composed of a (de-)parser and a sequence of processing stages. In such a switch, each stage is particularized to its usage. This specificity makes it impossible to trade memory size for another, add a new table, support new headers or new actions.




SDN wants multiple stages of match-action (flexible allocation), flexible actions and flexible header fields (OpenFlow was designed with these goals in mind). Alternative ways to enable flexibility don't compare in terms of cost-performance ratio: software is 100x too slow and expensive, NPU's are 10x too slow and expensive and FPGAs are 10x too slow and expensive.

The solution to all these problems are **Reconfigurable Match Tables (RMT)**. However, this is challenging to achieve. What kind of switch architecture could support flexibility and yet run at Terabits per second. At 1Tbps, packet size of 1000bits and 10 operations per packet, we'd need 10 billion operations per second. With a single processor, this would require a 10Ghz processor, which is not feasible. Parallelizing things with a packet-parallel architecture allows to have lower CPU speeds through duplication of the processing units. However, one issue is to scale the memory-to-CPU bandwidth. Replicating memory comes at a huge cost in die area. The solution is to organize the processing as a pipeline. Pipelined architectures organize processing through a sequence of processing units and local memory. For flexibility, each processing unit/memory can be made generic. Each CPU can process distinct packets, with up to 10 packets going through the pipeline simultaneously.



(a) Pipeline switch architecture



(b) Match/Action Forwarding Model

4.1 Parsing

Parsing is the (complex) process of identifying and extracting the appropriate fields in a packet header. Parsing faces multiple challenges:

- Throughput: Parser must run at line-rate (parse 1 packet every 70ns on a 10Gbps link)
- Dependency: Parsing involves sequential processing as headers typically point to the next
- Incompleteness: Some headers do not even identify the subsequent headers
- Heterogeneity: Many header formats exist that can appear in various orders/locations

Parse graphs are directed acyclic graphs encoding header types and their sequence. A parser can be divided into two separate blocks: header identification and field extraction. In a programmable parser, the two modules rely on runtime information (stored in memory, e.g. in RAM and/or TCAM) instead of hard-coded logic.

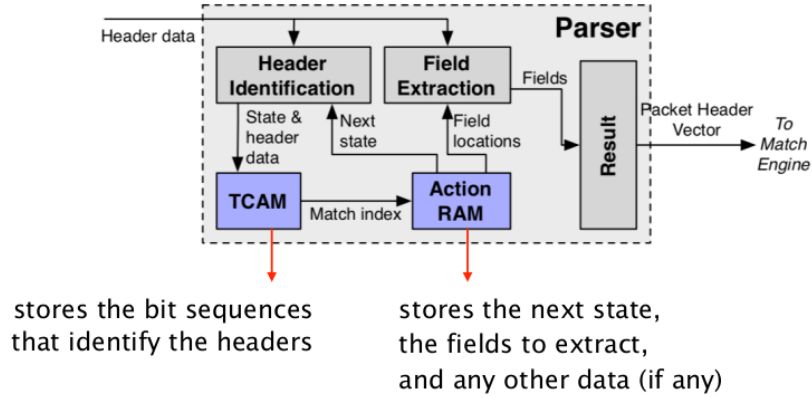


Figure 15: Programmable hardware parser

4.2 Logical pipeline

A compiler translates a given RMT logical pipeline (specified in P4) into a physical one. Each physical stage contains dedicated SRAM, for exact matches, and TCAM, for ternary matches. The compiler maps each individual logical stage to one or more physical stage. Small tables can share a stage (up to 16 per stage), while large tables can span multiple ones.

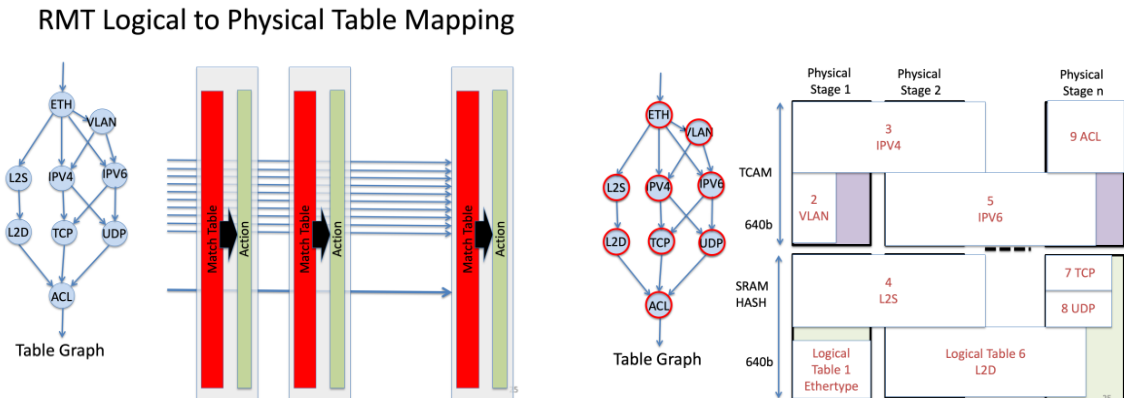


Figure 16: Logical pipeline [5]

The RMT pipeline relies on many Arithmetic Logic Units (ALU) to perform actions on the result of a match. Each ALU modifies only one word of a header (a header is composed of many words). Each stage of the RMT pipeline contains one ALU per word of the header vector (that's a lot of ALUs).

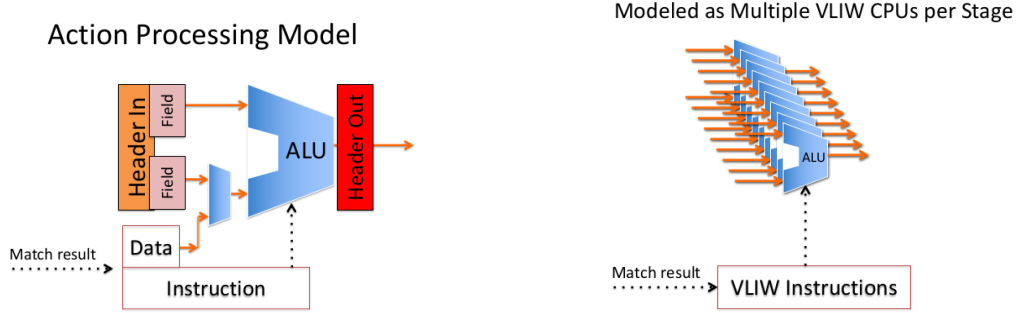


Figure 17: Logical pipeline [5]

Building a RMT pipeline is only 15% more expensive than building a fixed-function switching pipeline. The biggest cost is the memory not the processing logic.

In conclusion, we can design a flexible chip using the RMT switch model, bring processing close to the memories (pipeline of many stages) and bring processing to the wires (224 action CPUs per stage). This only causes a 15% cost increase.

4.3 Barefoot Tofino

Barefoot Tofino 6.5 Tbps backplane allows for several billion packets per second at line rate. Tofino relies on Packet Header Vector (PHV) to pass states between stages. Tofino uses four folded match-action pipeline in which the same stages are used for ingress and egress pipeline.

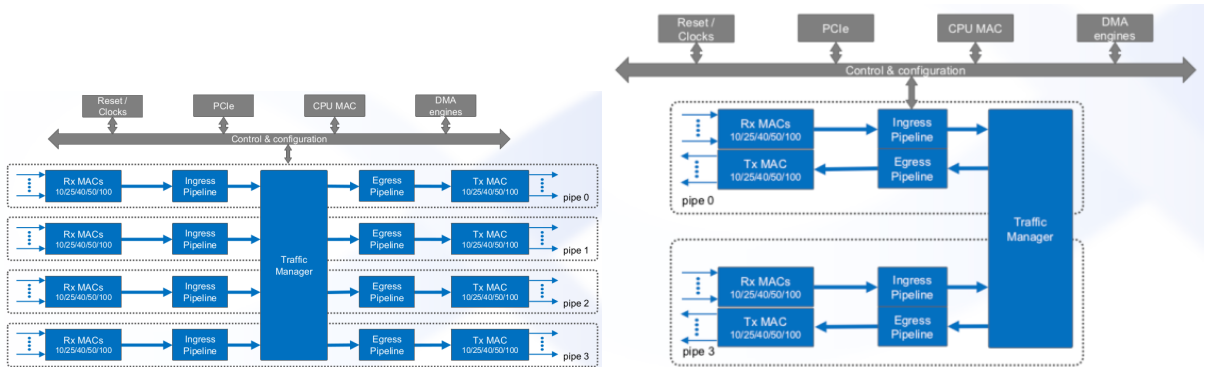



Figure 18: Barefoot Tofino pipeline [2]

4.3.1 Packet Header Vector (PHV)

PHVs are a set of uniform containers (8, 16, 32 bit) that carry the headers and metadata along the pipeline. Fields can be packed into any container or their combination. PHV allocation step in the compiler decides the actual packing.



4.3.2 Basic structure

Tofino basically follows the RMT pipeline




Figure 19: Barefoot Networks Tofino basic structure [2]

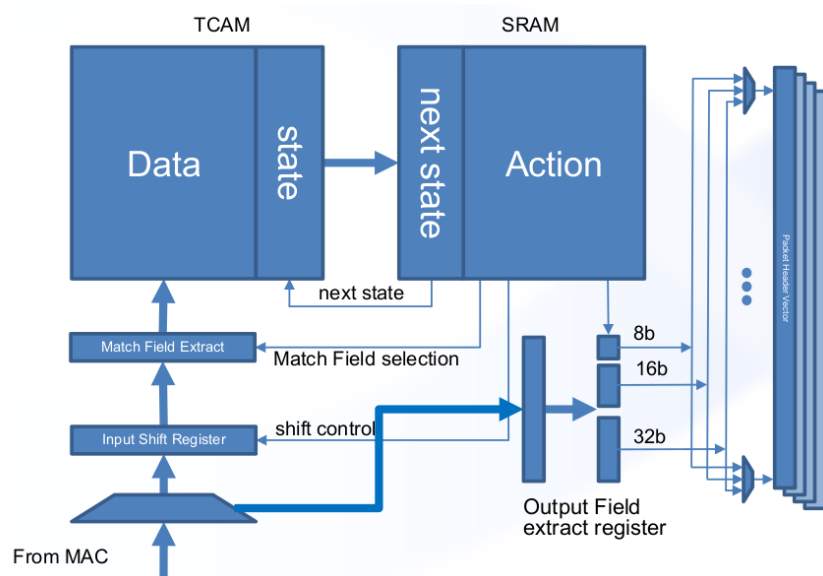


Figure 20: Barefoot Networks Tofino programmable parser [2]

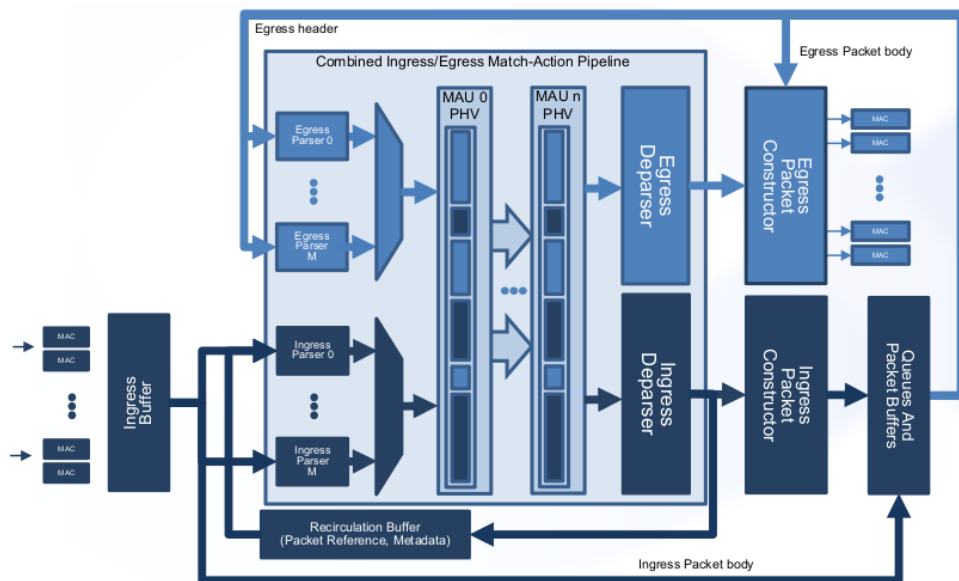


Figure 21: Barefoot Networks Tofino pipeline organization [2]

Each match action stage is regularly structured around: crossbars, memory units and ALUs.

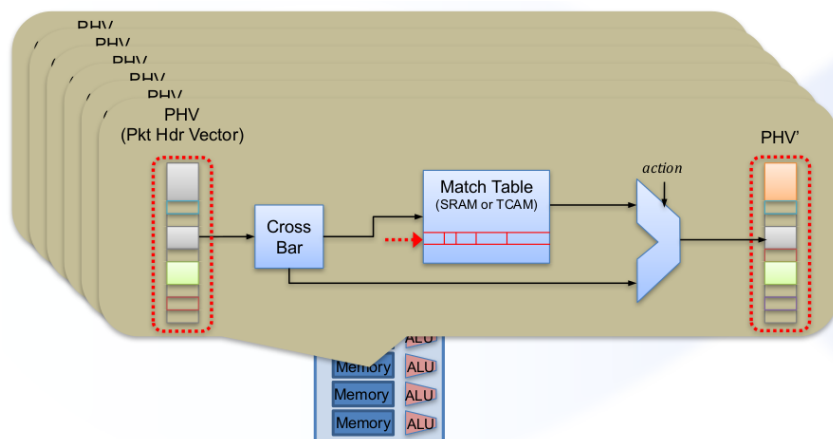


Figure 22: Tofino match action stage [2]

4.3.3 Parallel processing in Tofino

Most P4 programs have inherent parallelism. In Tofino, multiple tables allow multiple parallel lookups. All actions from all active tables are combined. Each PHV container has its own, independent processor.

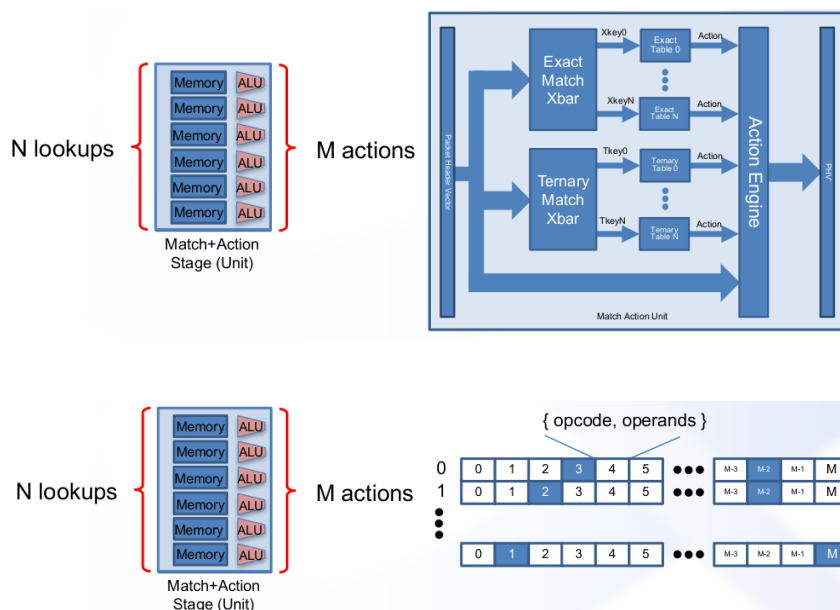


Figure 23: How Tofino supports parallelism [2]

4.4 P4 switch design

- 64 x 10Gb ports (960M packets/second, 1Ghz pipeline)
- Programmable parser
- 32 match/action stages
- Huge TCAM (10x current chips): 64k TCAM words x 640b
- SRAM hash tables for exact matches (128k words x 640b)
- 224 action processors per stage

5 P4-based applications

Current research in data plane programmability includes:

- Data plane programmability for performance, monitoring, applications offloading
- Platforms, correctness, management for data plane programmability

A few papers that I found interesting and insightful to read are:

- Fast network congestion detection and avoidance using P4
- Herding the Elephants: Detecting Network-Wide Heavy Hitters with Limited Resources
- Network-Wide Heavy Hitter Detection with Commodity Switches
- NetPaxos: Consensus at Network Speed

A large set of papers on programmable data planes aim at improving performance, esp. load balancing.

5.1 CONGA: Distributed Congestion-Aware Load Balancing for Datacenters

The main thesis of CONGA [1] is that datacenter load balancing is best done in the network instead of the transport layer, and requires global congestion-awareness to handle asymmetry. Unlike local schemes such as ECMP and Flare, CONGA seamlessly handles asymmetries in the topology or network traffic. CONGA leverages an existing datacenter overlay to implement a leaf-to-leaf feedback loop and can be deployed without any modifications to the TCP stack. CONGA senses the distant drumbeats of remote congestion and orchestrates flowlets to disperse evenly through the fabric. We leave to future work the task of designing more intricate rhythms for more general topologies.




Figure 24: CONGA: Distributed Congestion-Aware Load Balancing for Datacenters

5.2 INT: In-band Network Telemetry

INT is a mechanism for collecting network state in the dataplane:

- As close to realtime as possible
- At current and future line rates
- With a framework that can adapt over time

Examples of network state are: (switch id, ingress port id, egress port id), egress link utilization, hop latency, etc. P4 enables flexible packet parsing and modification for INT. P4 allows INT to adapt to any encapsulation format, any state required to be collected, any feature, protocol (current and future).

For example, HULA uses periodic INT probes to disseminate path utilization to switches. Further, flowlet detection and path selection happens at all switches which allows hop-by-hop adaptive routing and load balancing. HULA offers better scalability than CONGA.

In summary, INT provides real-time network state directly in the dataplane while scaling to arbitrarily large networks, scaling to current and future link speeds and being able to adapt to any network, any encap, any application. Knowledge of real-time network state opens up new possibilities such as enhanced monitoring and troubleshooting, network-state aware routing etc.

5.3 NetCache: Balancing Key-Value Stores with Fast In-Network Caching

NetCache [6] is a rack-scale key-value store that leverages in-network data plane caching to achieve billions QPS throughput with $\approx 10\mu s$ latency even under dynamic, skewed (i.e. few hot items) workload. It solves the problem of load-balancing in key-values stores observing dynamic, skewed workload. It leverages that a small but very fast cache can provide perfect load-balancing. NetCache relies on the $O(\text{billion})$ throughput of programmable network devices to achieve it in practice. It relies on a tailored UDP-based protocol, an de/encoding scheme for storing variable length values, and sketches. The challenge is to cache the hottest $O(N \log N)$ items with limited insertion rate. The goal is to react quickly and effectively to workload changes with minimal updates.

5.4 p4v: Practical Verification for Programmable Data Planes

How do you make sure that a programmable network works as it should? P4 is a low-level language \rightarrow many gotchas. Several problems can arise:

- header validity: it's hard to always keep validity in your head
- unambiguous forwarding: forwarding behavior depends on hardware

Challenges:

- imprecise semantics: P4 language spec doesn't give precise semantics
- modeling the control plane: Table rules are not statically known, populated by the control plane at runtime.

p4v [7] is an automated tool for verifying P4 programs. It considers all paths while still being practical for large programs.

5.5 Others

- LossRadar, FlowRadar: Develop techniques and tools to monitor all flows by relying on in-switch data structures (Bloom Filters) and decoding them at the controller-level.
- DAPPER, Network-wide HH: Develop P4-based detection mechanisms to diagnose TCP performance issue (e.g. small receiver buffers) heavy-hitter (e.g. port scanners, super-spreader, DDoS).
- SketchLearn, ElasticSketch, UnivMon: Introduce techniques to make sketch-based monitoring more practical (by making sketches adaptive or "universal")
- SwingState: manage programmable network: we need an OS for the data plane! SwingState is a state management framework with 1 primitive: moveStates.

6 Exercise design summaries

6.1 03-ECMP

Goal: implement ECMP, a very basic (but widely used) technique to load balance traffic across multiple equal cost paths.

6.1.1 Data plane

- Headers:
 - Standard ethernet, IP, TCP headers
 - metadata struct stores ecmp hash and ecmp group id
- Parser: Parse ethernet, IPv4, TCP
- Ingress:
 - `ipv4_lpm` MAT: matches on the dstIP (lpm). If the dstIP is either directly connected to the switch or there is only one path towards the dstIP, `set_nhop()` action is called to just rewrite egress port and MACs. If the dstIP is reachable over *multiple* paths, we call the `ecmp_group()` action with the ecmp group id and the number of next hops.
 - `ecmp_group_to_nhop` MAT: matches on the ECMP group id and ECMP group hash (exact). Calls `set_nhop()`. The rules specify which next hop to choose for which ecmp group id and ecmp hash.
 - `set_nhop()` action: rewrites MACs and egress port, decreases ipv4 TTL
 - `ecmp_group()` action: hashes the flow 5-tuple into `meta.ecmp_hash` and stores the ecmp group id in `meta.ecmp_group_id`.
 - apply: apply forwarding table. In case `ecmp_group` was called, apply `ecmp_group_to_nhop`
- Egress: -
- Checksum: Since we adapt the packet in `set_nhop` (decreasing the ttl), we need to recompute the ipv4 checksum.
- Deparser: emit all headers in order

6.2 03-Flowlet-switching

Goal: ECMP works well for flows with similar size. However, in real traffic flows vary vastly in size. ECMP suffers from performance problems if multiple large flows go through the same path due to hash collisions. We want to fix this by using information from `standard_metadata` to fix the collision problem of ECMP by implementing flowlet switching on top of ECMP.

6.2.1 Data plane

- Headers:
 - Standard ethernet, IP, TCP headers
 - metadata struct stores ecmp hash and ecmp group id, flowlet last timestamp, flowlet time difference, flowlet register index and flowlet id.
- Parser: Parse ethernet, IPv4, TCP
- Ingress:
 - Two registers: `flowlet_to_id` gives the flowlet id for the flowlet register index (hash of the flow 5-tuple). `flowlet_time_stamp` gives the last timestamp for the flowlet register index (hash of the flow 5-tuple).
 - `read_flowlet_registers` action: hashes the flow 5-tuple, reads the flowlet id and flowlet last timestamp into metadata, updates the timestamp with the current timestamp from `standard_metadata`.

- `update_flowlet_id` action: generate a random number and use it as the new flowlet id for the current flowlet (flowlet register index available in metadata)
- `ipv4_lpm` and `ecmp_group_to_nhop` MATs as in ECMP
- `set_nhop()` and `ecmp_group()` actions as in ECMP, however the hash in `ecmp_group()` now also includes the flowlet id (which periodically changes!)
- apply: If IPv4 header is valid, read flowlet registers and update flowlet id incase the flowlet time diff exceeds timeout. Then apply `ipv4_lpm` and `ecmp`.
- Egress: -
- Checksum: Since we adapt the packet in `set_nhop` (decreasing the ttl), we need to recompute the ipv4 checksum.
- Deparser: emit all headers in order

6.3 04-CountMin-sketch

Goal: Implement a CountMin sketch to estimate occurrences of distinct elements.

6.3.1 Data plane

- Headers:
 - Standard ethernet, IP, TCP headers
 - metadata struct stores all flow 5-tuple hashes and all counter values read from registers
- Parser: Parse ethernet, IPv4, TCP
- Ingress:
 - forwarding MAT: match on ingress port (exact) and rewrite egress port, using `set_egress_port()` action.
 - Define N registers for CountMin sketch
 - Define `sketch_count()` action that calls a macro for each register. The macro hashes the flow 5-tuple, reads the counter in the registers at the index, increases the counter by 1 and writes back to register.
 - apply: call `sketch_count()` if IPv4 and TCP headers are valid, then apply the forwarding MAT
- Egress: -
- Checksum: -
- Deparser: emit all headers in order

6.3.2 Control plane

- Write rules into forwarding MAT
- Read registers from control plane (`self.controller.register_read()`)
- Read CountMin sketch by hashing the flow and reading all registers at the specified indices. Return the minimum counter value.

6.4 04-Heavy-hitter-detector

Goal: Use counting bloom filter to estimate heavy hitters.

6.4.1 Data plane

- Headers:
 - Standard ethernet, IP, TCP headers
 - metadata struct stores two flow 5-tuple hashes and two counter values read from register
- Parser: Parse ethernet, IPv4, TCP
- Ingress:
 - ipv4_lpm MAT: match on dstIP (lpm) and rewrite MACs and egress port, using `ipv4_forward()` action.
 - `ipv4_forward()` action: takes dstMAC and egress port
 - Define 1 register for the counting bloom filter
 - Define `update_bloom_filter()` action that hashes the flow 5-tuple with two different hash functions, writes the hash values in metadata, reads the register at the hash indexes, store these values in metadata, increase counters and write back to register.
 - apply: if IPv4 and TCP headers are valid, call `update_bloom_filter()` and drop traffic and return in case both counters exceed a threshold. Apply `ipv4_lpm` at the end.
- Egress: -
- Checksum: Since we adapt the packet in `set_nhop` (decreasing the ttl), we need to recompute the ipv4 checksum.
- Deparser: emit all headers in order

6.5 05-Simple-routing

Goal: implement and provide a control plane to the ECMP routing exercise (see 6.1). Instead of specifying the entries for the forwarding tables manually, we will now implement a controller that generates and installs forwarding rules automatically, based on the network topology.

6.5.1 Data plane

Equivalent to 6.1.

6.5.2 Control plane

- Iterate over all switches in the topology and connect to them:

```
for p4switch in self.topo.get_p4switches():
    thrift_port = self.topo.get_thrift_port(p4switch)
    self.controllers[p4switch] = SimpleSwitchAPI(thrift_port)
```
- For each switch (source), iterate over all switch destinations. If the destination is equal to the source, install a `ipv4_lpm` table rule for all directly connected hosts. For all other destinations, check the shortest path from source to dst and iterate over all hosts connected to dst. If there is only one single shortest path, install a rule for it in `ipv4_lpm`. Otherwise, install a ECMP rule. ECMP groups are defined per switch with the set of next ports. A new ECMP group needs to be created in case the current selection of dst ports does not have an ECMP group yet.

6.6 05-Traceroutable

Goal: Extend a P4 router with the functionality to respond to traceroute packets, i.e. packets where the IPv4 TTL (time to live) value is equal to 1. If a router receives such a packet, it generates an ICMP Time Exceeded message and sends it back to the original sender of the expired packet.

Note: An ICMP header encapsulates the IP and TCP header of the original request.

6.6.1 Data plane

- Headers:
 - Standard ethernet, IP, TCP headers
 - new icmp header (8 bytes) with type, code, checksum, unused
 - headers struct contains icmp and two ipv4 headers (the original one and the one added by ICMP, `ipv4_icmp`). The `ipv4_icmp` header is going to be the ipv4 header of the ICMP packet with the actual src and dst. The original ipv4 header (ipv4 header of the packet that triggered the ICMP response) is going to be encapsulated in the ICMP header (unchanged).
 - metadata struct stores ecmp hash and ecmp group id
- Parser: Parse ethernet, IPv4, TCP
- Ingress:
 - `ipv4_lpm`, `ecmp_group_to_nhop` MATs as in ECMP
 - `set_nhop`, `ecmp_group` actions as in ECMP
 - `icmp_ingress_port` MAT: matches on the ingress port and calls action `set_src_icmp_ip`. The reason for this table is when sending back an ICMP reply, we need to set the srcIP of the ICMP packet. We set it to the IP of the interface on which we send out the packet. This table will be populated by the controller.
 - `set_src_icmp_ip`: takes an IP address as parameter and sets the `hdr.ipv4_icmp.srcAddr` to it.
 - `apply`: If the ipv4 header is valid and `ttl > 1`, we simply apply the `ipv4_lpm` table and do ECMP if the `ecmp_group` action was run (see 6.1). If ipv4 and tcp headers are valid *and* `ttl == 1`, we need to respond with an ICMP packet. We thus set the icmp and `ipv4_icmp` headers to valid, and prepare the packet to be sent back as ICMP packet: set egress port to ingress port, swap MACs, copy the original ipv4 header into the `ipv4_icmp` header, set its `dstAddr` to the `srcAddr` and apply the `icmp_ingress_port` MAT (i.e. setting the srcIP to the interface IP). Finally, set the protocol field of `ipv4_icmp` to ICMP, adapt the `ttl` and adapt the `totalLen`. Set the ICMP type and code and truncate the packet.
- Egress: -
- Checksum: since we have three different headers (`ipv4_icmp`, `icmp`, `ipv4`) that all include checksums and all headers can be changed, we have to recalculate the checksum for each header.
- Deparser: emit ethernet, `ipv4_icmp`, `icmp`, `ipv4`, `tcp`

6.6.2 Control plane

- Write rules into `ipv4_lpm` MAT
- Write ECMP rules as explained in section 6.5
- For each switch, for each interface on it, write an entry into `icmp_ingress_port` MAT.

```
for sw_name, controller in self.controllers.items():  
    for intf, node in self.topo.get_interfaces_to_node(sw_name).items():  
        ip = self.topo.node_to_node_interface_ip(sw_name, node).split("/")[0]  
        port_number = self.topo.interface_to_port(sw_name, intf)  
  
        self.controllers[sw_name].table_add("icmp_ingress_port",  
            "set_src_icmp_ip", [str(port_number)], [str(ip)])
```

6.7 06-Packet-Loss-Detection

Goal: Implement LossRadar. Packet losses are very common in nowadays networks. Those losses can be caused by many reasons: congestion, blackholes, faulty links or corrupted forwarding tables. LossRadar uses two IBLTs for each unidirectional link: one upstream meter for the sending switch, one downstream meter for the receiving switch. To make sure that the meters that are exchanged where exposed to the exact same set of packets batches are used. Lost packets can be detected by taking the difference between the UM and the DM, i.e. subtract the counts and xor the values. This will be done by the controller. The controller will be informed by a copy-to-cpu packet which includes a loss header with the last batch id.

6.7.1 Data plane

- Headers:
 - Standard ethernet, IP, TCP, UDP headers
 - loss header (batch_id, padding, nextProtocol), will be between IP and TCP/UDP header, thus we need a nextProtocol field
 - headers struct contains ethernet, ipv4, loss, tcp, udp
 - metadata struct stores tmp_src_port, tmp_dst_port (since we have TCP and UDP, we need uniform reference to the port), 3 hash value fields for the UM, 3 hash value fields for the DM, temp srcIP, dstIP, ports_proto_id, counter for updating the IBLT, fields with last local, previous and current batch id and an indicator whether to execute DM, UM.
- Parser: Parse ethernet, IPv4, loss, TCP/UDP
- Ingress:
 - forwarding MAT that calls `set_egress_port` action (set correct egress port)
 - `remove_loss_header` MAT that calls `remove_header` action (if the link goes toward a host we need to remove the loss header)
 - 8 registers: one each srcIP, dstIP, (ports, proto, id), counter; for both UM and DM
 - `compute_hash_indexes()` action: calculates hash indexes for all three hash functions on flow 5-tuple into metadata with correct bounds for both UM and DM.
 - `apply_um_meter()` action: increases counter and xors value in UM registers at all indexes indicated by hashes
 - `apply_dm_meter()` action: increases counter and xors value in DM registers at all indexes indicated by hashes




Figure 25: Register port and batch structure

- apply: get L4 ports in metadata, add loss header if not present (assumes that we have no host-switch-host connections) - else store batch id as previous batch id, read last local batch id and update batch id. If new batch id is different from last local batch id, inform controller. Update batch id in loss header. Compute hashes and remove loss header if necessary. Apply UM (if flag to not execute is not set), apply DM (if flag to not execute is not set). Finally, drop packet if ttl=1, otherwise decrease ttl by 1.
- Egress: if the packet is a clone packet, set loss header valid, ipv4 invalid, set batch id as last local batch id and set etherType.
- Checksum: Update checksum of ipv4
- Deparser: emit ethernet, ipv4, loss, tcp, udp

6.7.2 Control plane

- Write rules into forwarding MAT
- set crc custom hashes
- Upon receiving a clone packet:
 - Get the loss header fields (batch id) and the switch name (src)
 - Iterate over all neighboring switches (dst) of the switch (src)
 - For each neighboring switch, check the link to it: get the ports of the link at each respective switch, then extract the registers for both switches (UM for src, DM for dst).
 - Decode the two IBLTs: first, take the difference between UM and DM (UM-DM) (subtract counters and xor values) and store in diff array. Then, while there is a counter==1 in the counter diff array, get the array index associated with that counter and read out srcIP, dstIP, (ports, proto, id) in the respective diff arrays. From this, you get the flow 5-tuple. This allows you to calculate all 3 hashes for the flow. Then, at the 3 positions in the diff arrays indicated by the hashes, decrease the counter by 1 and xor the values with the values of the found flow 5-tuple. The found flow is a dropped packet.

7 Oral exam questions

What is P4? What does it enable?

P4 is a programming language that allows us to program reprogrammable switches. It tries to achieve reconfigurability, protocol independence and target independence. Rather than have the switch tell us the limited set of things it can do, P4 gives us a way to tell the switch what it should do, and how it should process packets. P4 enables us to define custom protocols by defining parsers and let's us control switches "top-down" by first specifying their forwarding behavior, then populating the tables we've defined. Instead of the switch chip vendor defining our API for us, and locking us into using their next chip as well, P4 let's us define the API we need to populate the switch. We say that P4 is "top-down" because it puts the network architect, programmer, and developer in charge, rather than the chip vendor. As such, we can define our own protocols and don't have to rely on hardware vendors to include protocols, as it was the case with OpenFlow. In OpenFlow, each supported protocol needs to be specified since OpenFlow assumes fixed-function switch chips.

What are the main differences between OpenFlow and P4?

OpenFlow is an API to the data plane, P4 allows to program the data plane.

P4 is *not* a replacement for OpenFlow - they try to achieve different things. Both are focused on opening up the forwarding plane. OpenFlow tries to centrally and remotely control lots of different switches, essentially providing an API to directly access the data plane. P4 on the other hand addresses a different need in the network, the need to *program* the data plane.

OpenFlow assumes the switch to have a fixed, well-known behavior, typically described in the datasheet of the switch ASIC. These switches supported a fixed set of protocols and you couldn't change their behavior or add new protocols. As interest in OpenFlow grew, more and more header-types were added. However, adding your own protocol is essentially only possible if you are a company with big influence and smaller companies are bound to use existing protocols. Adding more and more protocols led to OpenFlow becoming an "overblown beast". This strange behavior is all caused by the fact that switches were not programmable.

That's where P4 comes in. With programmable switches, we can simply tell the switch how to process packets and what tables to keep. Programmers could define whatever API made sense to populate the tables they created in the switches. And that exactly is what P4 does. P4 is a common language to specify how a switch should process packets. P4 lets us define what headers a switch will recognize (or "parse"), how to match on each header, and what actions we would like the switch to perform on each header.

What roles do tables, actions and control flow play in P4?

Tables, actions and control flow are all components of P4 that play together for packet processing. P4 is based on the reconfigurable match tables (RMT) model. Once a packet has been parsed, it is matched against a set of tables.

- **Tables:** A table matches on a set of header fields and applies an action with specified parameters. A match-action table entry thus consists of a set of header fields to match and a set of action parameters. Matching can be done in different ways (exact, lpm, ternary, range).
- **Action:** An action is essentially a *function* as known from classical programming languages. Actions can modify the packets and can be either invoked from within the control block or as a result of a match in a match-action table.
- **Control flow:** Control flow allows us to specify how a packet is processed, specifying in what order tables and actions are applied. We can run actions, apply a table, check if a table has matched, which action was run, etc.

What different memory structures exist in P4 switches? How do they differ?

We know that P4 switches make packet processing more efficient by making use of pipelined architectures that organize processing through a sequence of processing units and local memory. Match-action tables and parsers both need efficient memory lookups for matching. However, different matching types require different kinds of memory. As such, P4 switches have two different types of memory: TCAM and SRAM.

- **SRAM(input: address, output: data):** SRAM allows for fast random access, however it is expensive. As such, SRAM is used for exact matching in match action tables and to store the next state, fields to extract and other data (if any) for parsing.
 - **Exact match MAT:** In an exact match, we want to lookup the action for exact match-key data. We can use Cuckoo hashing to get $O(1)$ worst-case lookup, deletion and average-case insertion. For exact matching we hash the match-key and then do an SRAM lookup.
 - **Parsing:** The SRAM stores next state information, fields to extract, and any other data needed during parsing.
- **TCAM(input: match field value, output: address of action):** TCAM (ternary content-addressable memory) is a specialized type of high-speed memory that searches its entire contents in a single clock cycle. Generally speaking, CAM is often described as the opposite of random access memory (RAM). To retrieve data on RAM, the operating system (OS) must provide the memory address where the data is stored. Data stored on CAM can be accessed by performing a query for the content itself, and the memory retrieves the addresses where that data can be found. Due to its parallel nature, CAM (and by extension TCAM) is much faster than RAM. CAM is essentially a massively parallel lookup engine that searches all entries in parallel and selects the best match in constant time. The term ternary refers to the memory's ability to store and query data using three different inputs: 0, 1 and X. The X input, which is often referred to as a don't care or wildcard state, enables TCAM to perform broader searches based on pattern matching (a binary CAM matches every bit precisely).
 - **LPM/ternary match MAT:** In an LPM match, we want to lookup the action with the most specific match for a match-key. As such, SRAM is not ideal since we have multiple possible memory locations (a key can match several patterns). This is where TCAM comes into play: TCAM matches the key against all entries and returns the action address of the best match.
 - **Parsing:** The TCAM stores bit sequences that identify headers. The current state and a subset of bytes from the buffer are sent to the TCAM, which returns the

first matching entry. The output of the TCAM points to an SRAM entry which specifies the next state for the header identification module and the headers that were matched. The RAM entry may also specify data such as the number of bytes to advance and the subset of bytes to extract next (see figure 15).

What types of constructs do we need to maintain state across packets?

In P4, variables have local scope and their value is not maintained across subsequent invocations for different packets. However, there is a need for stateful programming (e.g. think of a stateful firewall).

P4 offers 4 different types of stateful objects: registers, counters and meters.

- **Registers:** Registers are arrays of bit vectors. Both the width of the bit vector and the length of the array can be specified. Registers are stateful - they maintain their value across packets. We can read from and write to a specified register index. Registers can also be read from the control plane.
- **Counters:** Counters are used for counting. Counters are assigned in arrays and an index needs to be specified when executing a `count()` operation. Counters can be increased from the data plane but can only be read from the control plane. Counters can be one of three different types: packets, bytes, packet_and_bytes. There also exist direct counters which can be directly attached to a MAT.
- **Meters:** Meters are used for rate-limiting. Meters are also assigned in arrays. Like counters, we also have direct meters that can be attached to MATs.
- **MATs:** Match-action tables can also be used to maintain state.

What is a bloom filter? How do you dimension one?

A bloom filter is a probabilistic data structure that is used for membership queries. A Bloom filter enables a deterministic number of required operations at the cost of false positives. A Bloom filter consists of a fixed-size array (size M) of 1-bit cells and K hash functions for insertion of N elements. Inserting an element consists of hashing the element with all hash functions and setting the 1-bit cell to 1 at each position indicated by the K hash functions. To check if an element is in the Bloom filter, we hash the element with each Hash function and check the 1-bit cells indexed by the hashes. If all cells are 1, we declare the element to be in the filter, else it's not in the filter. Due to hash collisions, we can get false positives, i.e. in the case where element e_1 has a hash collision with element e_2 on all K hash functions:

$$h_i(e_1) = h_i(e_2) \quad \forall i \in [1, K]$$

Given N, M , we can design a Bloom Filter with minimal false positive rate by choosing

$$K = \ln(2) * \frac{M}{N}$$

(see section 3.1.1)

Explain what is a counter and the different ways it can be used.

Counters are one out of three stateful objects in P4. Counters are used for counting (obviously) and are always assigned in arrays. Counters can be one of three different counting types: packets, bytes, packets_and_bytes. Further, there are two different counter types: indirect and direct counters. We can call `count()` from the data plane but counters can only be read from the control plane.

- **Indirect counter:** Indirect counters are arrays of `n_counters` independent counter values. Calling `count(in S index)` will increase the counter at the specified index.
- **Direct counter:** Direct counters are counters that are directly attached to tables.

Describe the role of the 'register' extern in the v1model and how it can be used (at a high level).

Registers are one out of three stateful objects in P4. Registers are useful for storing (small amounts of) arbitrary data. Registers are assigned in arrays for which we can specify the length and width. We can read and write from/to registers from the data plane as well as from the control plane. This sets them apart from counters which can only be read from the control plane. As such, registers are widely used to store data that must be available over subsequent packets.

Further, registers play an important role when creating more complex data structures such as Bloom filters or CountMin sketches.

In the exercise we made a 'traceroutable' P4 network. To do that we needed to generate ICMP packets when the TTL expired. Explain which additions were needed? (header definitions, how did the main headers struct look like, parser, main ingress/ egress logic, checksum updates, etc).

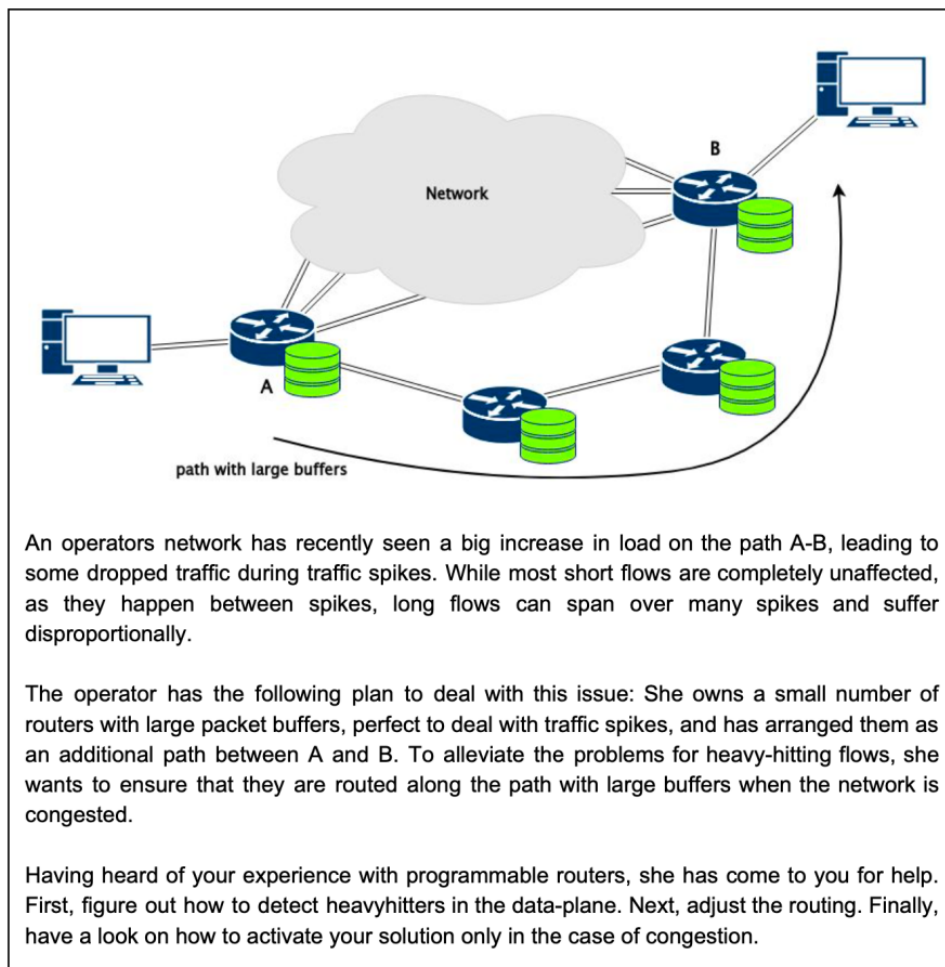
ICMP packets consist of an ethernet header, IP header and ICMP header. Inside the ICMP header, we have the original IP and TCP headers of the packet that created the ICMP message.

- Headers: ICMP header (8 bytes total: type, code, checksum, unused), `ipv4_icmp` headers which is a normal `ipv4` header.
- Main headers struct: the main headers struct thus included the following structs: `ethernet` (ethernet header), `ipv4_icmp` (`ipv4` header), `icmp` (icmp header), `ipv4` (`ipv4` header), `tcp` (tcp header)

The parser remained unchanged since we only sent out ICMP packets. In the Deparser, we had to add two additional emits, essentially emitting all headers in the right order: `ethernet`, `ipv4_icmp`, `icmp`, `ipv4`, `tcp`.

For the ingress, we needed one additional table which matches on the ingress port and then sets the `srcIP` of the `ipv4_icmp` to the IP of the interface. In the ingress logic, we have to differentiate between `ttl > 1`, in which case we apply normal forwarding/ECMP, and `ttl == 1`. For the later, we first set the `ipv4_icmp` and `icmp` headers to valid. We then adapt the packet headers such that the packet can be sent back out on the port it came in (i.e. switch MACs, set egress port to ingress port, set `dstIP` to `srcIP`). We then apply the aforementioned table to set the `srcIP` of the `ipv4_icmp` header to the IP of the interface.

Design question



There are three parts to this system design that need to be specified: the heavy hitter detection, the routing change and how to only activate the solution in case of congestion.

- Heavy hitter detection: all flows enter the network through a single switch A. Thus no distributed detection technique is necessary. We can thus use a CountMin sketch to estimate the number of packets in a flow. The `count()` action is called for every valid packet in the ingress.
- Routing change: We can use two different `ipv4_lpm` MATs: a 'normal' MAT `ipv4_lpm` that routes traffic through the standard network and a buffer MAT `ipv4_lpm_buffer` that send incoming traffic at A over the large buffer path. Both MATs match on the `dstIP` (`lpm`) and the action sets the next hop by rewriting the egress port and the `dstMAC`. The next hops of the `ipv4_lpm` MAT are inside the network whereas the next hop of the `ipv4_lpm_buffer` MAT is the next large buffer switch after A. In the apply section of the ingress, we first increase the CountMin sketch and then compare the counts (stored in a metadata field) against a pre-specified threshold. If all counts exceed the threshold, we have a heavy hitter and apply the `ipv4_lpm_buffer` MAT which will set the next hop to a large buffer router. Else we apply the `ipv4_lpm` MAT. This approach does the routing change fully in the data plane, thus there is no induced latency by control plane communication. However, we do need to maintain two MATs which wastes memory.
- Activation: In order to only activate the solution in case of congestion, we need to observe some congestion signals to actually detect congestion.

References

- [1] ALIZADEH, M., EDSALL, T., DHARMAPURIKAR, S., VAIDYANATHAN, R., CHU, K., FINGERHUT, A., LAM, V. T., MATUS, F., PAN, R., YADAV, N., AND ET AL. Conga: Distributed congestion-aware load balancing for datacenters. *SIGCOMM Comput. Commun. Rev.* 44, 4 (Aug. 2014), 503514.
- [2] BAREFOOT NETWORKS. Barefoot networks: Programmable data plane at terabit speeds, 2020.
- [3] BOSSHART, P., DALY, D., GIBB, G., IZZARD, M., MCKEOWN, N., REXFORD, J., SCHLESINGER, C., TALAYCO, D., VAHDAT, A., VARGHESE, G., ET AL. P4: Programming protocol-independent packet processors. *ACM SIGCOMM Computer Communication Review* 44, 3 (2014), 87–95.
- [4] BOSSHART, P., GIBB, G., KIM, H.-S., VARGHESE, G., MCKEOWN, N., IZZARD, M., MUJICA, F., AND HOROWITZ, M. Forwarding metamorphosis: Fast programmable match-action processing in hardware for sdn. *SIGCOMM Comput. Commun. Rev.* 43, 4 (Aug. 2013), 99110.
- [5] GIBB, G., VARGHESE, G., HOROWITZ, M., AND MCKEOWN, N. Design principles for packet parsers. In *Architectures for Networking and Communications Systems* (2013), IEEE, pp. 13–24.
- [6] JIN, X., LI, X., ZHANG, H., SOULÉ, R., LEE, J., FOSTER, N., KIM, C., AND STOICA, I. Netcache: Balancing key-value stores with fast in-network caching. In *Proceedings of the 26th Symposium on Operating Systems Principles* (New York, NY, USA, 2017), SOSP 17, Association for Computing Machinery, p. 121136.
- [7] LIU, J., HALLAHAN, W., SCHLESINGER, C., SHARIF, M., LEE, J., SOULÉ, R., WANG, H., CAUNDEFINEDCAVAL, C., MCKEOWN, N., AND FOSTER, N. P4v: Practical verification for programmable data planes. In *Proceedings of the 2018 Conference of the ACM Special Interest Group on Data Communication* (New York, NY, USA, 2018), SIGCOMM 18, Association for Computing Machinery, p. 490503.
- [8] NETWORKED SYSTEMS GROUP (NSG). Advanced topics in communication networks course website and slides, 2020.
- [9] THE P4 LANGUAGE CONSORTIUM. P4(16) language specification, 2020.